

# Assignment 3 - Using pretrained CNNs for image classification

---

## Original Assignment Description

In the previous assignments involving classification, we were performing a kind of simple feature extraction on images by making them greyscale and flattening them to a single vector. This vector of pixel values was then used as the input for some kind of classification model.

For this assignment, we're going to be working with an interesting kind of cultural phenomenon - fashion. On UCloud, you have access to a dataset of *Indo fashion* taken from this [Kaggle dataset](#). There is a paper which goes along with it on *arXiv.org*, which you can read [here](#).

Your instructions for this assignment are short and simple:

- You should write code which trains a classifier on this dataset using a *pretrained CNN like VGG16*
- Save the training and validation history plots
- Save the classification report

## Tips

- You should not upload the data to your repo - it's around 3GB in size.
  - Instead, you should document in the README file where your data comes from, how a user should find it, and where it should be saved in order for your code to work correctly.
- The data comes already split into training, test, and validation datasets. You can use these in a **TensorFlow** data generator pipeline like we saw in class this week - you can see an example of that [here](#).
- There are a lot of images, around 106k in total. Make sure to reserve enough time for running your code!
- The image labels are in the metadata folder, stored as JSON files. These can be read into **pandas** using `read_json()`. You can find the documentation for that online.

## Assignment 3 - Luke Ring (202009983)

---

### Contribution

The code for this assignment was written independently and is my own (Luke Ring, 202009983) [zeyus @ github](#).

### Description

This repository contains code that fine tunes a VGG16 model on the Indo Fashion dataset. The training history plot, confusion matrix and classification report are saved in the **out/** folder, along with a csv file containing the training history, and a txt file containing the model summary and classification report.

The data were preprocessed using the following steps:

- Proportionally rescaled the images to `IMAGE_SIZExIMAGE_SIZE` pixels, or `IMAGE_WIDTHxIMAGE_HEIGHT` pixels if specified.
- Images were zero-padded to match the input size of the model, with the image centered.
- Normalized the pixel values to be between 0 and 1.
- For training data, random horizontal flips were and random rotations of -20 to 20 percent were applied.

For the results below, the images were resized to 100x200 pixels, and the batch size was 256.

## Setup

### Windows GPU (optional)

If you have an NVIDIA GPU you can do the following before installing the prerequisites:

- Install [Anaconda](#)
- Create a new environment using `conda create -n vgg16 python=3.9`
- Activate the environment using `conda activate vgg16`
- Install cudatoolkit and cudnn with `conda install -c conda-forge cudatoolkit=11.2 cudnn=8.1.0`

### Prerequisites

- Install the required packages using `pip install -r requirements.txt`
- Download the dataset from [Kaggle](#), unzip and save it to `data/` either manually, or by doing the following:
  - Get your kaggle API token from [here](#) and save it to `~/.kaggle/kaggle.json`
  - Run `python src/cnn.py --download` to download the dataset

## Usage

- Run `python src/cnn.py` to train the model with the default parameters.

Most settings can be customized, such as input resizing, batch size, number of epochs, etc. Run `python src/cnn.py --help` to see all the available options.

```
> python .\src\cnn.py --help
usage: cnn.py [-h] [--version] [-m MODEL_SAVE_PATH] [--download] [-d DATASET_PATH]
[-s IMAGE_SIZE] [-w IMAGE_WIDTH] [-t IMAGE_HEIGHT] [-b BATCH_SIZE] [-e EPOCHS] [-o
OUT] [-n] [-c FROM_CHECKPOINT] [-r] [-p PARALLEL]
```

Text classification CLI

optional arguments:

```
-h, --help                show this help message and exit
--version                show program's version number and exit
-m MODEL_SAVE_PATH, --model-save-path MODEL_SAVE_PATH
                        Path to save the trained model(s) (default: models)
--download                Download the dataset from kaggle (default: False)
-d DATASET_PATH, --dataset-path DATASET_PATH
                        Path to the dataset (default: data)
```

```

-s IMAGE_SIZE, --image-size IMAGE_SIZE
                        The image size (width and height) (default: 32)
-b BATCH_SIZE, --batch-size BATCH_SIZE
                        The batch size (default: 32)
-e EPOCHS, --epochs EPOCHS
                        The number of epochs (default: 10)
-o OUT, --out OUT      The output path for the plots and stats (default: out)
-n, --no-train         Do not train the model (default: False)
-c FROM_CHECKPOINT, --from-checkpoint FROM_CHECKPOINT
                        Use the checkpoint at the given path (default: None)
-r, --resnet           Use ResNet50 as the base model. (default: False)
-p PARALLEL, --parallel PARALLEL
                        Number of workers/threads for processing. (default: 4)

```

## Results

The final VGG16 model was trained in two steps, with a batch size of 256 and 150 epochs each.


The classification report, model history and confusion matrix are described below.

### First training step


#### Classification report

precision	recall	f1-score	support		
	blouse	0.93	0.89	0.91	500
	dhoti_pants	0.84	0.50	0.63	500
	dupattas	0.74	0.54	0.62	500
	gowns	0.71	0.45	0.55	500
	kurta_men	0.67	0.81	0.73	500
	leggings_and_salwars	0.59	0.74	0.66	500
	lehenga	0.86	0.78	0.82	500
	mojaris_men	0.86	0.84	0.85	500
	mojaris_women	0.84	0.86	0.85	500
	nehru_jackets	0.90	0.81	0.85	500
	palazzos	0.88	0.63	0.74	500
	petticoats	0.64	0.86	0.73	500
	saree	0.60	0.91	0.72	500
	sherwanis	0.84	0.70	0.76	500
	women_kurta	0.53	0.74	0.62	500
	accuracy			0.74	7500
	macro avg	0.76	0.74	0.74	7500
	weighted avg	0.76	0.74	0.74	7500

#### Model history



Confusion matrix



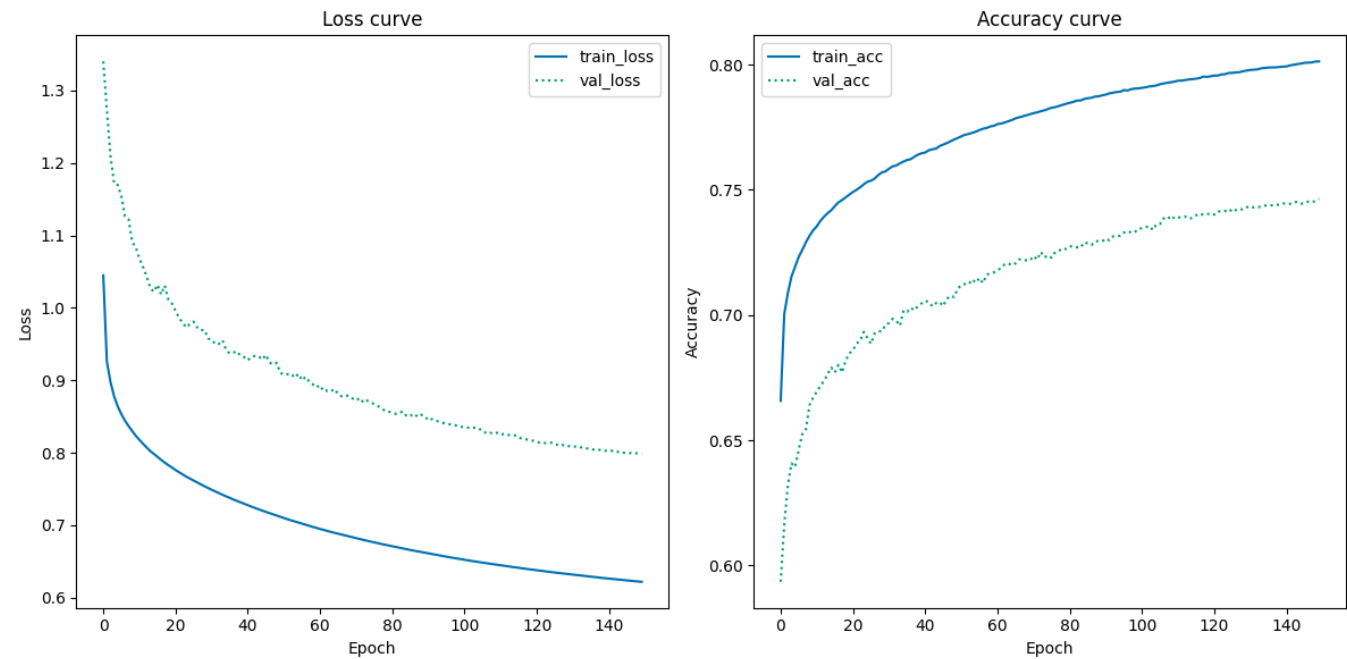
Second training step

Classification report


precision	recall	f1-score	support		
	blouse	0.90	0.94	0.92	500
	dhoti_pants	0.85	0.55	0.67	500
	dupattas	0.79	0.47	0.59	500
	gowns	0.69	0.41	0.52	500
	kurta_men	0.66	0.87	0.75	500
leggings_and_salwars		0.66	0.75	0.70	500
	lehenga	0.90	0.81	0.86	500
	mojaris_men	0.89	0.79	0.84	500
	mojaris_women	0.81	0.90	0.85	500

nehru_jackets	0.90	0.80	0.85	500
palazzos	0.89	0.68	0.77	500
petticoats	0.75	0.83	0.79	500
saree	0.71	0.91	0.80	500
sherwanis	0.87	0.68	0.77	500
women_kurta	0.45	0.88	0.60	500
accuracy			0.75	7500
macro avg	0.78	0.75	0.75	7500
weighted avg	0.78	0.75	0.75	7500

Model history



Confusion matrix



## Conclusion

The model was able to achieve an accuracy of 75% on the test set, which is not bad, but it could have been better. Notably, something must have gone wrong with the training continuation, although the fine-tuned model was loaded correctly, the model training did not seem to continue, but given more epochs it's possible that the accuracy could have been further improved.